

Heart Disease Prediction using Machine Learning Algorithms

Import Libraries

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from matplotlib import rcParams
from matplotlib.cm import rainbow
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
from sklearn import *

from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
```

Read dataset and preprocess the data

```
In [ ]: heartData = pd.read_csv("./dataset/heart.csv")
```

```
In [ ]: heartData.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null    int64
1   sex         303 non-null    int64
2   cp          303 non-null    int64
3   trestbps    303 non-null    int64
4   chol        303 non-null    int64
5   fbs         303 non-null    int64
6   restecg     303 non-null    int64
7   thalach     303 non-null    int64
8   exang       303 non-null    int64
9   oldpeak     303 non-null    float64
10  slope       303 non-null    int64
11  ca          303 non-null    int64
12  thal        303 non-null    int64
13  target      303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

```
In [ ]: heartData.describe()
```

```
Out[ ]:
```

	age	sex	cp	trestbps	chol	fbs
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000
75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000

Total missing % of data

```
In [ ]: missing_data = heartData.isnull().sum()
total_percentage = (missing_data.sum()/heartData.shape[0])*100

print(f"Total percentage of missing data is {round(total_percentage, 2)}%")

duplicate = heartData[heartData.duplicated()]
print("Duplicate rows: ")
duplicate

#drop duplicate rows
heartData = heartData.drop_duplicates()
```


Total percentage of missing data is 0.0%

Duplicate rows:

```
In [ ]: rcParams['figure.figsize'] = 10, 10

plt.matshow(heartData.corr())
plt.yticks(np.arange(heartData.shape[1]), heartData.columns)
plt.xticks(np.arange(heartData.shape[1]), heartData.columns)
plt.colorbar()
```

Out[]: <matplotlib.colorbar.Colorbar at 0x3034f8290>



Divide the data into training and testing classes

```
In [ ]: X = heartData.drop(['target'], axis=1)
y = heartData['target']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.
```

```
print("XTrain: ",X_train.shape[0],"XTest: ",X_test.shape[0])
print("Ytrain: ",y_train.shape[0],"Ytest: ",y_test.shape[0])
```

```
XTrain: 211 XTest: 91
Ytrain: 211 Ytest: 91
```

Model Building

KNN Algorithm

```
In [ ]: knn_scores = []

for k in range(2, 21):
    knn_classifier = KNeighborsClassifier(n_neighbors = k)
    knn_classifier.fit(X_train.values, y_train.values)
    knn_score = round(knn_classifier.score(X_test.values, y_test.values),
    knn_scores.append(knn_score)

knn_classifier = KNeighborsClassifier(n_neighbors=5)
knn_classifier.fit(X_train, y_train)
knn_score = knn_classifier.predict(X_test)
print(classification_report(y_test, knn_score))
```

	precision	recall	f1-score	support
0	0.61	0.49	0.54	45
1	0.58	0.70	0.63	46
accuracy			0.59	91
macro avg	0.60	0.59	0.59	91
weighted avg	0.60	0.59	0.59	91

```
In [ ]: score_knn = round(accuracy_score(knn_score, y_test)*100,2)
print(f"The accuracy score achieved using KNN is {score_knn}")
```

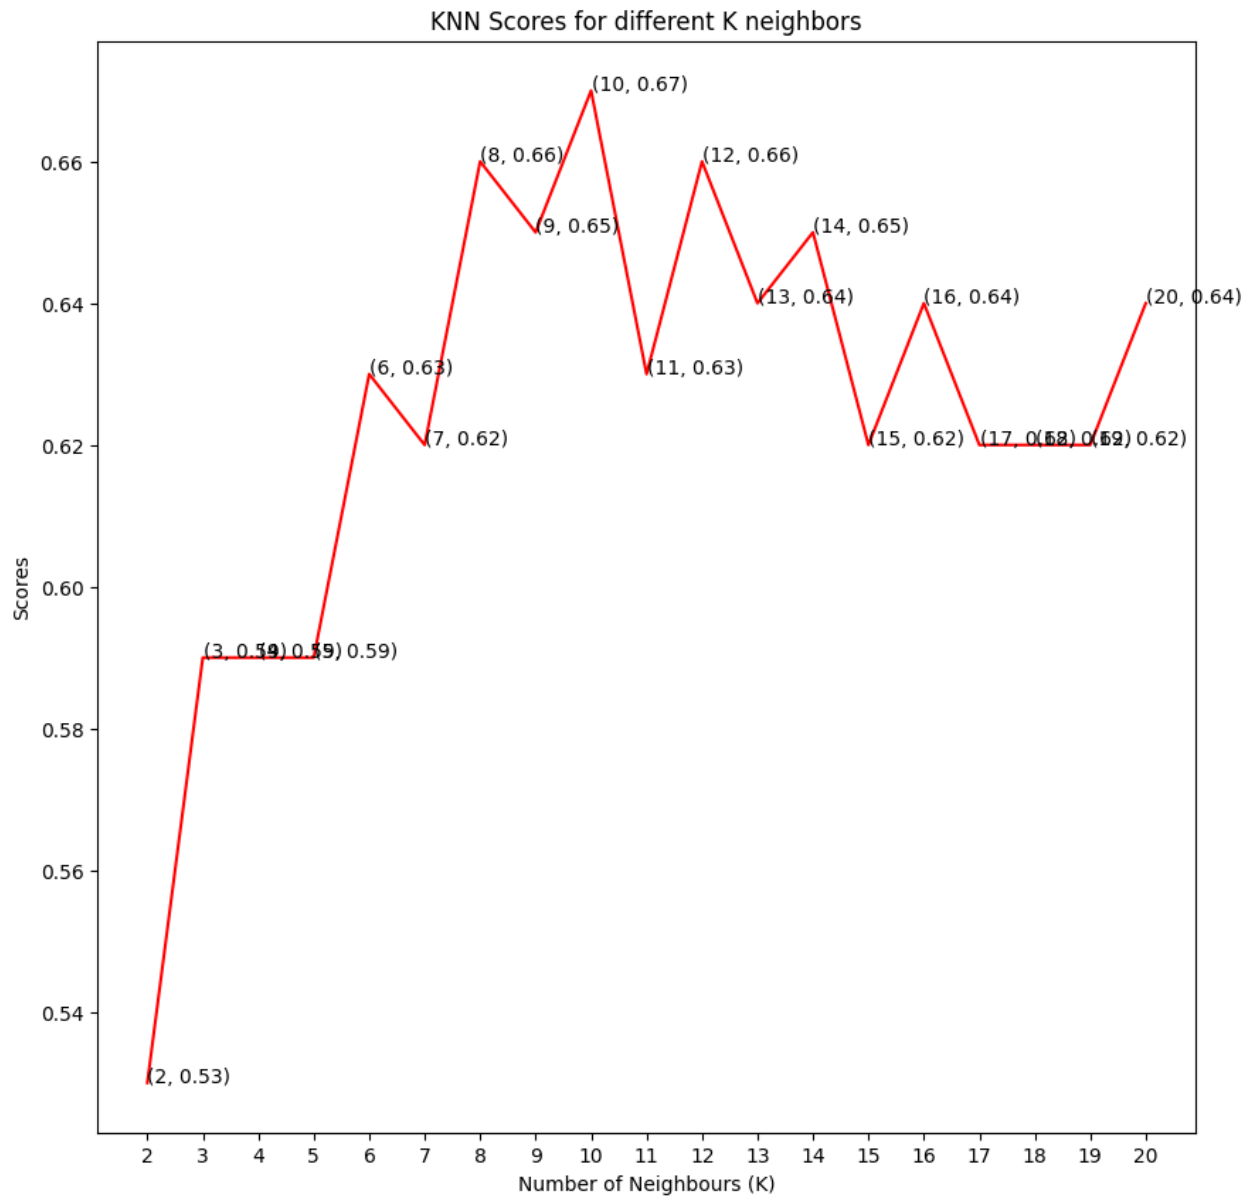
The accuracy score achieved using KNN is 59.34

```
In [ ]: plt.plot([k for k in range(2, 21)], knn_scores, color='red')

for i in range(2,21):
    plt.text(i, knn_scores[i-2], (i, knn_scores[i-2]))

plt.xticks([i for i in range(2,21)])
plt.xlabel("Number of Neighbours (K)")
plt.ylabel('Scores')
plt.title('KNN Scores for different K neighbors')
```

```
Out[ ]: Text(0.5, 1.0, 'KNN Scores for different K neighbors')
```



Support Vector Machine

```
In [ ]: from sklearn.metrics import accuracy_score

svc_scores = []
kernels = ['linear', 'poly', 'rbf', 'sigmoid']

for i in range(len(kernels)):
    svc_classifier = SVC(kernel=kernels[i])
    svc_classifier.fit(X_train.values, y_train.values)
    svc_scores.append(round(svc_classifier.score(X_test.values, y_test.values), 2))

svc_classifier = SVC(kernel = kernels[0])
svc_classifier.fit(X_train.values, y_train.values)
svc_prediction_result = svc_classifier.predict(X_test.values)

# print(svc_prediction_result)
```

```
print("Accuracy Score: ", accuracy_score(y_test.values, svc_prediction_re
```


Accuracy Score: 0.8241758241758241

```
In [ ]: colors = rainbow(np.linspace(0, 1, len(kernels)))

plt.bar(kernels, svc_scores, color=colors)

for i in range(len(kernels)):
    plt.text(i, svc_scores[i], svc_scores[i])
plt.xlabel('Kernels')
plt.ylabel('Scores')
plt.title('SVM scores Activation function wise')
```

Out[]: Text(0.5, 1.0, 'SVM scores Activation function wise')



```
In [ ]: score_svm = round(accuracy_score(svc_prediction_result,y_test)*100,2)
print(f"The accuracy score achieved using Linear SVM is {score_svm}")
```

The accuracy score achieved using Linear SVM is 82.42

Decision Tree

```
In [ ]: dt_scores = []

for i in range(1, len(X.columns)+1):
    dt_classifier = DecisionTreeClassifier(max_features = i, random_state=0)
    dt_classifier.fit(X_train.values, y_train.values)
    dt_scores.append(round(dt_classifier.score(X_test.values, y_test.values), 2))

print("Dt_scores: ", dt_scores)
```

Dt_scores: [0.67, 0.7, 0.79, 0.78, 0.79, 0.76, 0.76, 0.84, 0.73, 0.77, 0.76, 0.73, 0.74]

```
In [ ]: score_dt = round(max(dt_scores)*100, 2)
print(f"The accuracy score achieved using Decision Tree is {score_dt}%")
```

The accuracy score achieved using Decision Tree is 84.0%

```
In [ ]: dt_classifier = DecisionTreeClassifier(max_features=13, random_state=0)
dt_classifier.fit(X_train.values, y_train.values)
```

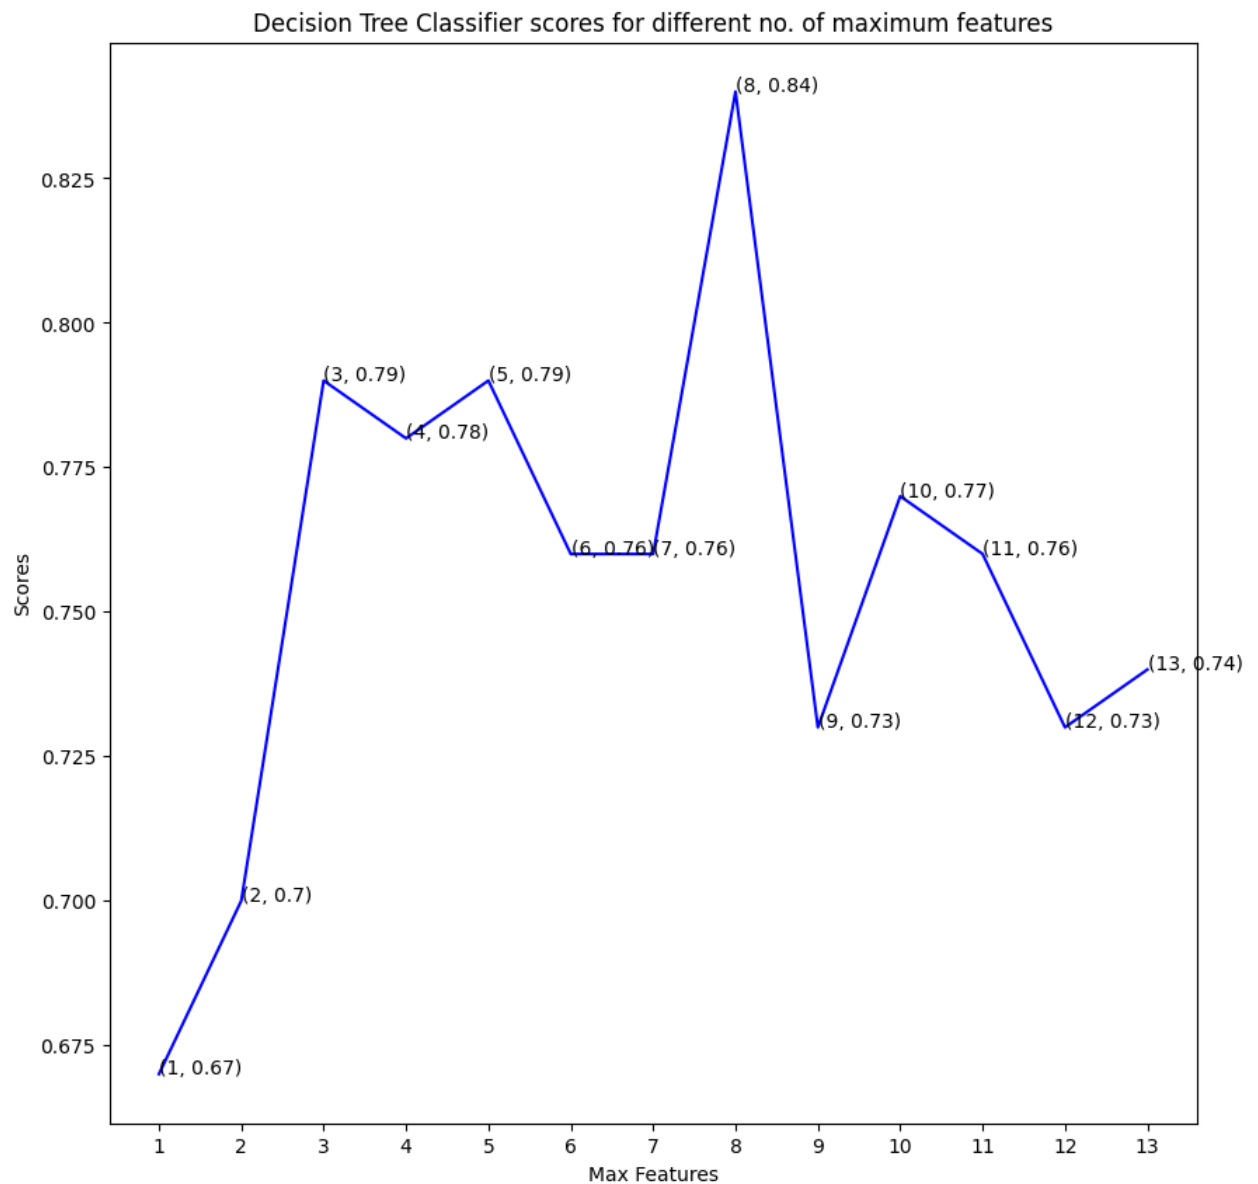
```
Out [ ]: ▼ DecisionTreeClassifier
DecisionTreeClassifier(max_features=13, random_state=0)
```

```
In [ ]: plt.plot([i for i in range(1, len(X.columns)+1)], dt_scores, color='blue')

for i in range(1, len(X.columns)+1):
    plt.text(i, dt_scores[i-1], (i, dt_scores[i-1]))

plt.xticks([i for i in range(1, len(X.columns)+1)])
plt.xlabel('Max Features')
plt.ylabel('Scores')
plt.title('Decision Tree Classifier scores for different no. of maximum f
```

```
Out [ ]: Text(0.5, 1.0, 'Decision Tree Classifier scores for different no. of max
imum features')
```



Random Forest

```
In [ ]: from sklearn.ensemble import RandomForestClassifier
```

```
In [ ]: rf_scores = []

estimators = [10, 20, 100, 200, 500]

for i in estimators:
    rf_classifier = RandomForestClassifier(n_estimators = i, random_state=42)
    rf_classifier.fit(X_train.values, y_train.values)
    rf_scores.append(round(rf_classifier.score(X_test.values, y_test.values), 2))
```

```
In [ ]: colors = rainbow(np.linspace(0, 1, len(estimators)))

plt.bar([i for i in range(len(estimators))], rf_scores, color = colors, width=0.5)


for i in range(len(estimators)):
```



```
plt.text(i, rf_scores[i], rf_scores[i])

plt.xlabel('No. of estimators')
plt.ylabel('Scores')
plt.title('Random Forest Classifier scores for different no. of estimator
```

Out[]: Text(0.5, 1.0, 'Random Forest Classifier scores for different no. of estimators')



```
In [ ]: score_rf = round(max(rf_scores)*100,2)
print(f"The accuracy score achieved using Random Forest is {score_rf}")
```

The accuracy score achieved using Random Forest is 86.0

Logistic Regression

```
In [ ]: from sklearn.linear_model import LogisticRegression
```

```
In [ ]: logistic_model = LogisticRegression()
logistic_model.fit(X_train.values, y_train.values)
logistic_model_prediction = logistic_model.predict(X_test.values)

print("Accuracy score: ", accuracy_score(y_test.values, logistic_model_prediction))
print(classification_report(y_test.values, logistic_model_prediction))
```

```
Accuracy score: 0.8131868131868132
```

	precision	recall	f1-score	support
0	0.87	0.73	0.80	45
1	0.77	0.89	0.83	46
accuracy			0.81	91
macro avg	0.82	0.81	0.81	91
weighted avg	0.82	0.81	0.81	91

```
In [ ]: score_lr = round(accuracy_score(logistic_model_prediction, y_test)*100,2)
print(f"The accuracy scored achieved using Logistic Regression is {score_lr}")
```

The accuracy scored achieved using Logistic Regression is 81.32

Naive Bayes

```
In [ ]: from sklearn.naive_bayes import GaussianNB
```

```
In [ ]: nb = GaussianNB()
nb.fit(X_train, y_train)

Y_pred_nb = nb.predict(X_test)
```

```
In [ ]: score_nb = round(accuracy_score(Y_pred_nb, y_test)*100,2)
print(f"The accuracy score achieved using naive bayes is: {score_nb}%")
```

The accuracy score achieved using naive bayes is: 82.42%

Output Final Score

```
In [ ]: scores = [score_lr, score_svm, score_knn, score_dt, score_rf, score_nb]
algorithms = ["Logistic Regression", "Support Vector Machine", "K-Nearest Neighbors", "Decision Tree", "Random Forest", "Naive Bayes"]

for i in range(len(algorithms)):
    print(f"{algorithms[i]} Accuracy: {scores[i]}%")
```

Logistic Regression Accuracy: 81.32%
 Support Vector Machine Accuracy: 82.42%
 K-Nearest Neighbors Accuracy: 59.34%
 Decision Tree Accuracy: 84.0%
 Random Forest Accuracy: 86.0%
 Naive Bayes Accuracy: 82.42%

Save Trained Models

```
In [ ]: import pickle
```

Saving the logistic regression model

```
In [ ]: pickle.dump(logistic_model, open('logistic_regression.model', 'wb'))
```

Saving the SVM Model

```
In [ ]: pickle.dump(svc_classifier, open('svm.model', 'wb'))
```

Saving Decision Tree Classifier

```
In [ ]: pickle.dump(dt_classifier, open('dt.model', 'wb'))
```

Saving Naive Bayes Classifier

```
In [ ]: pickle.dump(nb, open('naive_bayes.model', 'wb'))
```

Saving the KNN Model

```
In [ ]: pickle.dump(knn_classifier, open('knn.model', 'wb'))
```